

```
cd C:\Users\dhant\OneDrive\Desktop\simplilearn\Capstone\Project 2\
Healthcare - Diabetes
```

```
C:\Users\dhant\OneDrive\Desktop\simplilearn\Capstone\Project 2\
Healthcare - Diabetes
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import pylab as p
import missingno as msno
import warnings
warnings.filterwarnings('ignore')

data=pd.read_csv('health care diabetes.csv')
data.head(3)
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI |
|---|-------------|---------|---------------|---------------|---------|------|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |

1. Understand the dataset:

a. Identify the shape of the dataset

```
data.shape
```

```
(768, 9)
```

b. Identify the size of the dataset

```
data.size
```

```
6912
```

c. Identify the columns of the dataset

```
data.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',  
      'Insulin',
```

```
    'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
    dtype='object')
```

d. Identify the data types of the dataset

```
data.dtypes
```

```
Pregnancies      int64  
Glucose           int64  
BloodPressure     int64  
SkinThickness     int64  
Insulin           int64  
BMI               float64  
DiabetesPedigreeFunction float64  
Age               int64  
Outcome           int64  
dtype: object
```

e. Identify the information of the dataset

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

| # | Column | Non-Null Count | Dtype |
|---|--------------------------|----------------|---------|
| 0 | Pregnancies | 768 non-null | int64 |
| 1 | Glucose | 768 non-null | int64 |
| 2 | BloodPressure | 768 non-null | int64 |
| 3 | SkinThickness | 768 non-null | int64 |
| 4 | Insulin | 768 non-null | int64 |
| 5 | BMI | 768 non-null | float64 |
| 6 | DiabetesPedigreeFunction | 768 non-null | float64 |
| 7 | Age | 768 non-null | int64 |
| 8 | Outcome | 768 non-null | int64 |

```
dtypes: float64(2), int64(7)
```

```
memory usage: 54.1 KB
```

f. identifying the number of unique values of dataset

```
data.nunique()
```

```
Pregnancies      17  
Glucose           136  
BloodPressure     47  
SkinThickness     51  
Insulin           186  
BMI               248  
DiabetesPedigreeFunction 517  
Age               52  
Outcome           2  
dtype: int64
```

Information of data:

a. Identifying the total profile report of dataset

```
import pandas_profiling as pp
from pandas_profiling import ProfileReport
pp.ProfileReport(data)
```

```
{"version_major":2,"version_minor":0,"model_id":"dda5798e932c463eb270f
ef44af83e9c"}
```

```
{"version_major":2,"version_minor":0,"model_id":"b261d0051a3646eb9130c
fe4107addb9"}
```

```
{"version_major":2,"version_minor":0,"model_id":"5b53fa6f70ba47e7bf970
c7ee412ce1b"}
```


<IPython.core.display.HTML object>

skewness of data and its visualization

```
print(data.skew() )
p.plot(data.skew())
print( '\nSkewness for data : ')
```

| | |
|--------------------------|-----------|
| Pregnancies | 0.901674 |
| Glucose | 0.173754 |
| BloodPressure | -1.843608 |
| SkinThickness | 0.109372 |
| Insulin | 2.272251 |
| BMI | -0.428982 |
| DiabetesPedigreeFunction | 1.919911 |
| Age | 1.129597 |
| Outcome | 0.635017 |
| dtype: | float64 |

Skewness for data :

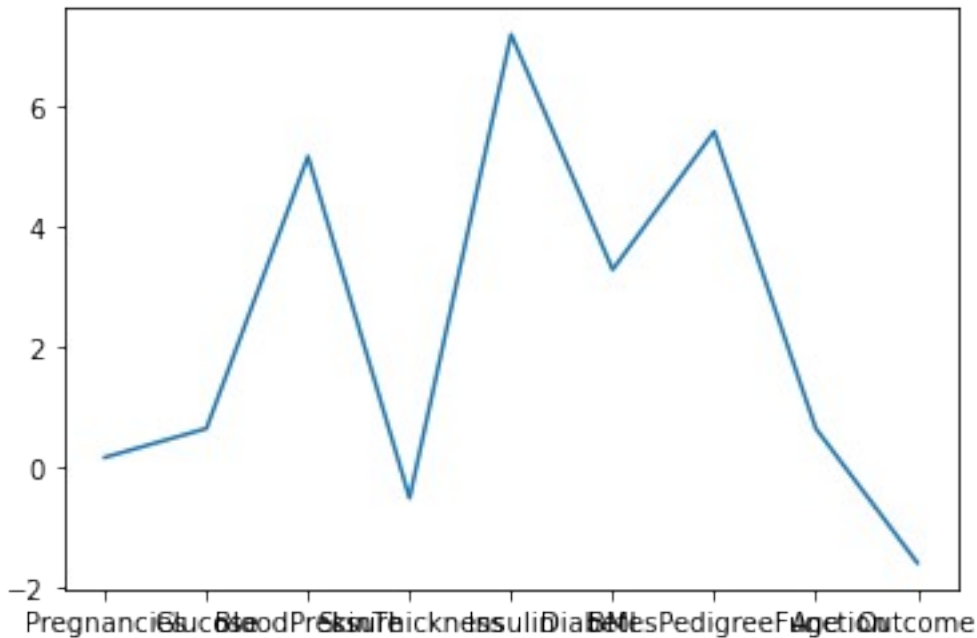


kurtosis and its visualization

```
print(data.kurtosis())
p.plot(data.kurtosis())
print( '\nkurtosis of data : ')
```

```
Pregnancies      0.159220
Glucose          0.640780
BloodPressure    5.180157
SkinThickness   -0.520072
Insulin          7.214260
BMI              3.290443
DiabetesPedigreeFunction  5.594954
Age              0.643159
Outcome         -1.600930
dtype: float64
```

kurtosis of data :



Statistical summary

```
data.describe().style.background_gradient(axis=1,cmap=sns.light_palette('green', as_cmap=True))
```

```
<pandas.io.formats.style.Styler at 0x1e045ff5940>
```


checking missing values

```
print(data.isnull().sum())
print("Display the missing values : there is no missing values","\n")
msno.bar(data)
plt.show()
```

```
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
```

```
dtype: int64
```

```
Display the missing values : there is no missing values
```



```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

| # | Column | Non-Null Count | Dtype |
|---|--------------------------|----------------|---------|
| 0 | Pregnancies | 768 non-null | int64 |
| 1 | Glucose | 768 non-null | int64 |
| 2 | BloodPressure | 768 non-null | int64 |
| 3 | SkinThickness | 768 non-null | int64 |
| 4 | Insulin | 768 non-null | int64 |
| 5 | BMI | 768 non-null | float64 |
| 6 | DiabetesPedigreeFunction | 768 non-null | float64 |
| 7 | Age | 768 non-null | int64 |
| 8 | Outcome | 768 non-null | int64 |

```
dtypes: float64(2), int64(7)
```

```
memory usage: 54.1 KB
```

Data Preprocessing

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

| # | Column | Non-Null Count | Dtype |
|---|---------------|----------------|-------|
| 0 | Pregnancies | 768 non-null | int64 |
| 1 | Glucose | 768 non-null | int64 |
| 2 | BloodPressure | 768 non-null | int64 |
| 3 | SkinThickness | 768 non-null | int64 |

```

4   Insulin          768 non-null   int64
5   BMI              768 non-null   float64
6   DiabetesPedigreeFunction  768 non-null   float64
7   Age              768 non-null   int64
8   Outcome          768 non-null   int64

```

```
dtypes: float64(2), int64(7)
```

```
memory usage: 54.1 KB
```

By observing the data file there are some variable with zero(0) that means the values is missing values...so the values replace with nan and treat by technique

```

data[["Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI"]]
= data[["Glucose", "BloodPressure", "SkinThickness", "Insulin",
"BMl"]].replace(0, np.NaN)

```

```
data.head()
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | |
|-------|-------------|---------|---------------|---------------|---------|------|
| BMI \ | | | | | | |
| 0 | 6 | 148.0 | 72.0 | 35.0 | NaN | 33.6 |
| 1 | 1 | 85.0 | 66.0 | 29.0 | NaN | 26.6 |
| 2 | 8 | 183.0 | 64.0 | NaN | NaN | 23.3 |
| 3 | 1 | 89.0 | 66.0 | 23.0 | 94.0 | 28.1 |
| 4 | 0 | 137.0 | 40.0 | 35.0 | 168.0 | 43.1 |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |
| 3 | 0.167 | 21 | 0 |
| 4 | 2.288 | 33 | 1 |


```

print(data.isnull().sum())
print("Display the missing values : there is some missing values","\n")
msno.bar(data)
plt.show()

```

| | |
|--------------------------|-----|
| Pregnancies | 0 |
| Glucose | 5 |
| BloodPressure | 35 |
| SkinThickness | 227 |
| Insulin | 374 |
| BMI | 11 |
| DiabetesPedigreeFunction | 0 |
| Age | 0 |

```
Outcome
dtype: int64
Display the missing values : there is some missing values
```



Observations

treating missing values with mean

Replacing NaN with mean values


```
data["Glucose"].fillna(data["Glucose"].mean(), inplace = True)
data["BloodPressure"].fillna(data["BloodPressure"].mean(), inplace = True)
data["SkinThickness"].fillna(data["SkinThickness"].mean(), inplace = True)
data["Insulin"].fillna(data["Insulin"].mean(), inplace = True)
data["BMI"].fillna(data["BMI"].mean(), inplace = True)
```

```
print(data.isnull().sum())
print("After treating the missing values : there is no missing values", "\n")
msno.bar(data)
plt.show()
```

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age              0
Outcome           0
```


dtype: int64

After treating the missing values : there is no missing values



```
data.describe().style.background_gradient(axis=1,cmap=sns.light_palette('green', as_cmap=True))
```

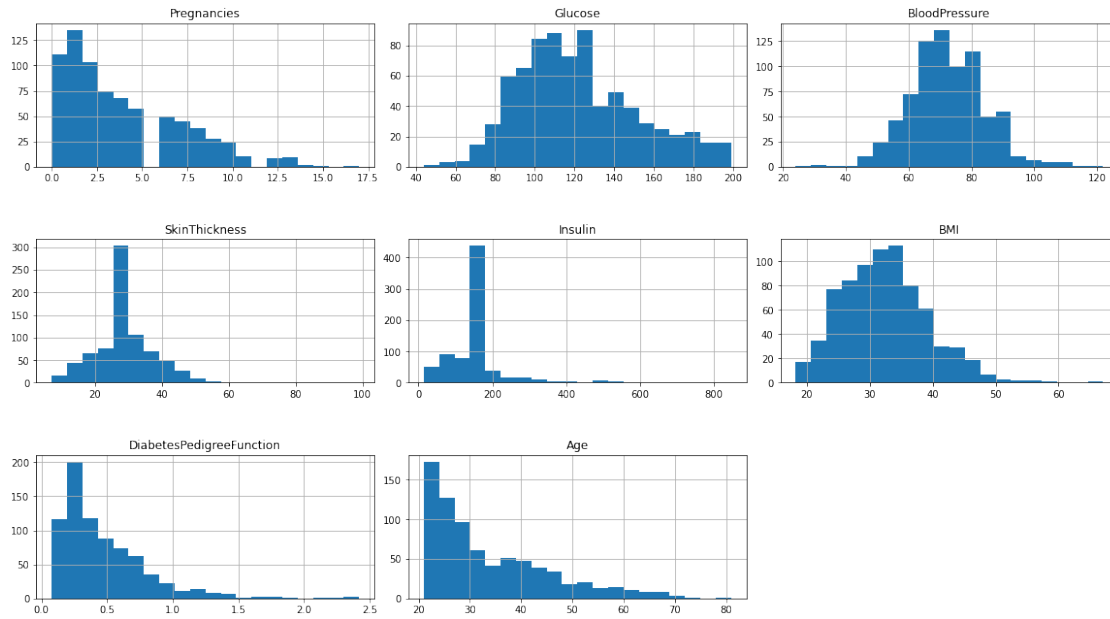
<pandas.io.formats.style.Styler at 0x2e0e5ccd340>

There are integer and float data type variables in this dataset. Create a count (frequency) plot describing the data types and the count of variables.


```
import itertools
```

```
col = data.columns[:8]
plt.subplots(figsize = (20, 15))
length = len(col)
```

```
for i, j in itertools.zip_longest(col, range(length)):
    plt.subplot((length/2), 3, j + 1)
    plt.subplots_adjust(wspace = 0.1, hspace = 0.5)
    data[i].hist(bins = 20)
    plt.title(i)
plt.show()
```



```
size=[10,20,52,45,85,65,41,25,35]
plt.axis("equal")
plt.pie(size,labels=data.columns,autopct="%1.1f%
%",radius=3,shadow=True,explode=[0.1,0.1,0.1,0.1,0.1,0.1,0.1,0.1,0.1],
startangle=60,counterlock=False)
plt.show()
```



Project Task: Week 2

1. Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of action.


```
print(data['Outcome'].value_counts())
print('The total number of outcomes of 0 are 500 and the number of
outcomes of 1 are 268')
sns.countplot('Outcome', data=data)
```

```
0    500
1    268
```

```
Name: Outcome, dtype: int64
```

```
The total number of outcomes of 0 are 500 and the number of outcomes
of 1 are 268
```

```
<AxesSubplot:xlabel='Outcome', ylabel='count'>
```

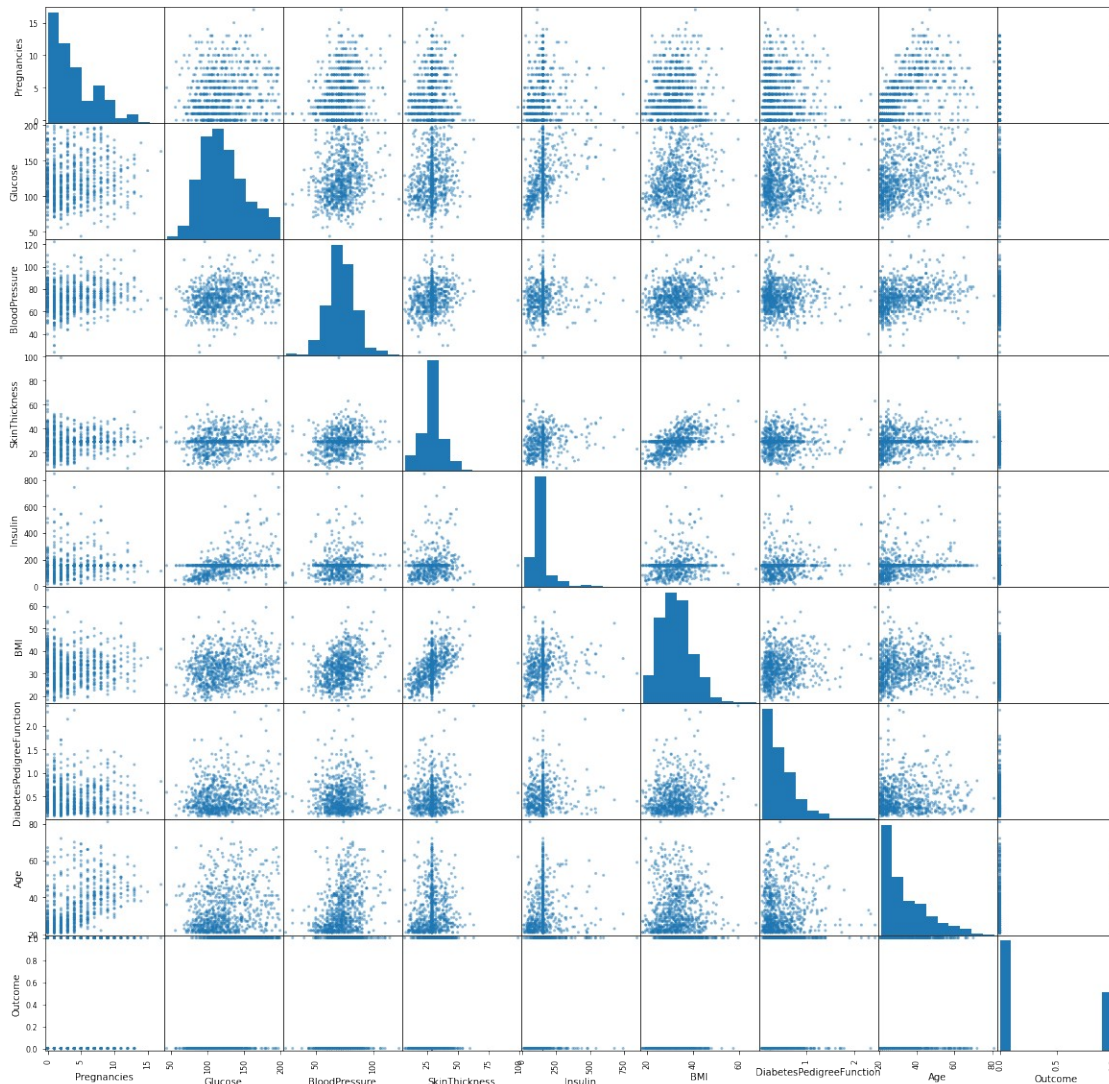


2. Create scatter charts between the pair of variables to understand the relationships.

Describe your findings.

Scatter plot matrix

```
from pandas.plotting import scatter_matrix  
scatter_matrix(data, figsize = (20, 20));
```



3. Perform correlation analysis. Visually explore it using a heat map.

```
print('correlation of each coulms along with that values')
data.corr()
```

correlation of each coulms along with that values

| | Pregnancies | Glucose | BloodPressure |
|-----------------|-------------|----------|---------------|
| SkinThickness \ | | | |
| Pregnancies | 1.000000 | 0.127911 | 0.208522 |
| 0.082989 | | | |
| Glucose | 0.127911 | 1.000000 | 0.218367 |
| 0.192991 | | | |
| BloodPressure | 0.208522 | 0.218367 | 1.000000 |
| 0.192816 | | | |
| SkinThickness | 0.082989 | 0.192991 | 0.192816 |
| 1.000000 | | | |
| Insulin | 0.056027 | 0.420157 | 0.072517 |
| 0.158139 | | | |


| | | | |
|--------------------------|-----------|----------|-----------|
| BMI | 0.021565 | 0.230941 | 0.281268 |
| 0.542398 | | | |
| DiabetesPedigreeFunction | -0.033523 | 0.137060 | -0.002763 |
| 0.100966 | | | |
| Age | 0.544341 | 0.266534 | 0.324595 |
| 0.127872 | | | |
| Outcome | 0.221898 | 0.492928 | 0.166074 |
| 0.215299 | | | |

| | Insulin | BMI | DiabetesPedigreeFunction |
|--------------------------|----------|----------|--------------------------|
| \ | | | |
| Pregnancies | 0.056027 | 0.021565 | -0.033523 |
| Glucose | 0.420157 | 0.230941 | 0.137060 |
| BloodPressure | 0.072517 | 0.281268 | -0.002763 |
| SkinThickness | 0.158139 | 0.542398 | 0.100966 |
| Insulin | 1.000000 | 0.166586 | 0.098634 |
| BMI | 0.166586 | 1.000000 | 0.153400 |
| DiabetesPedigreeFunction | 0.098634 | 0.153400 | 1.000000 |
| Age | 0.136734 | 0.025519 | 0.033561 |
| Outcome | 0.214411 | 0.311924 | 0.173844 |

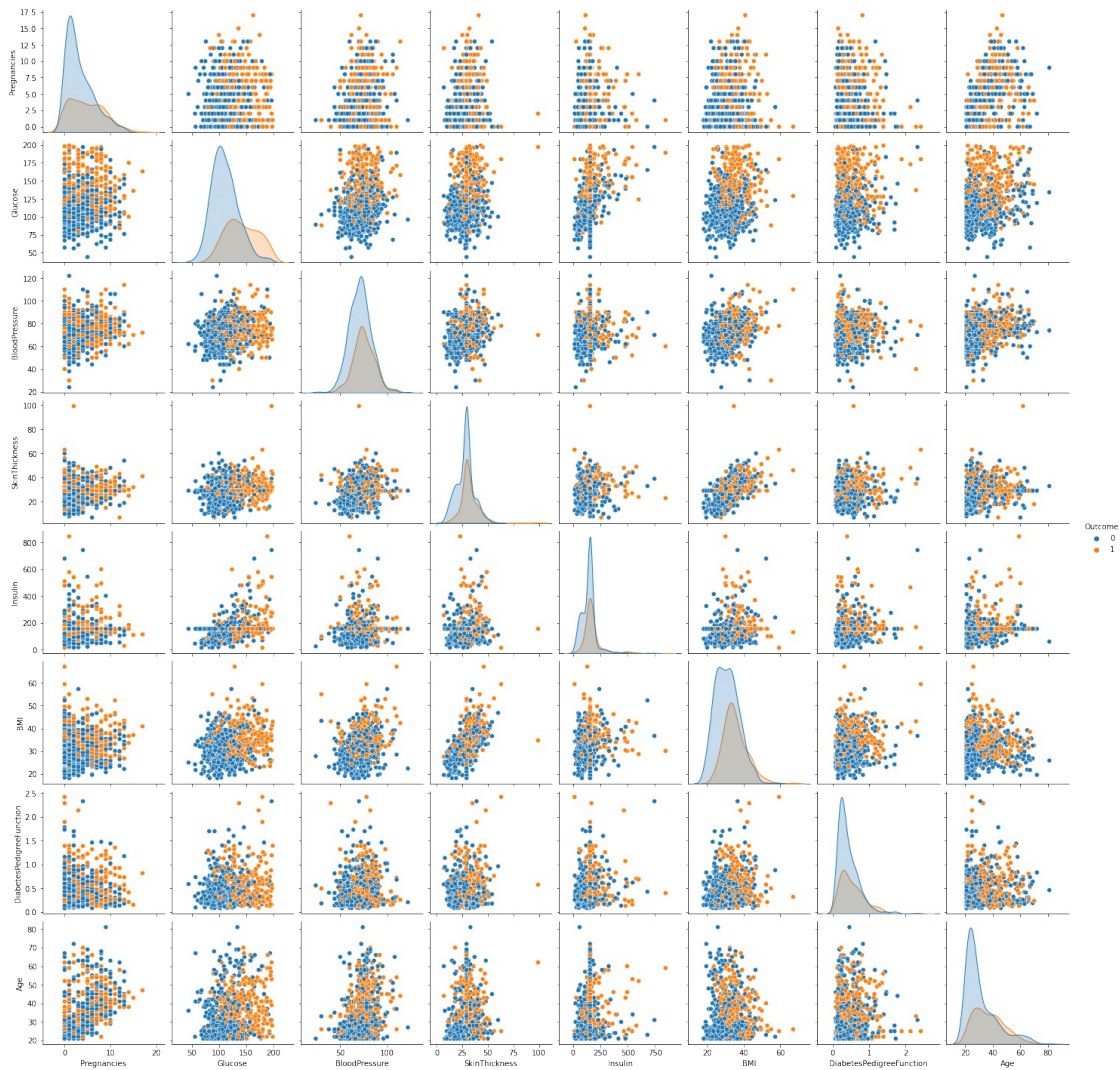
| | Age | Outcome |
|--------------------------|----------|----------|
| Pregnancies | 0.544341 | 0.221898 |
| Glucose | 0.266534 | 0.492928 |
| BloodPressure | 0.324595 | 0.166074 |
| SkinThickness | 0.127872 | 0.215299 |
| Insulin | 0.136734 | 0.214411 |
| BMI | 0.025519 | 0.311924 |
| DiabetesPedigreeFunction | 0.033561 | 0.173844 |
| Age | 1.000000 | 0.238356 |
| Outcome | 0.238356 | 1.000000 |

3. Plot the heatmap

```
plt.figure(figsize=(10,10))
heat_map = sns.heatmap( data.corr(), linewidth = 1 , annot = True)
plt.title( "HeatMap using Seaborn Method" )
plt.show()
```



```
# Pairplot
sns.pairplot(data = data, hue = 'Outcome')
plt.show()
```



Observations:

Project Task: Week 3

1. Devise strategies for model building. It is important to decide the right validation framework. Express your thought process.

`data.head(3)`

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin |
|-------|-------------|---------|---------------|---------------|------------|
| BMI \ | | | | | |
| 0 | 6 | 148.0 | 72.0 | 35.00000 | 155.548223 |
| 33.6 | | | | | |
| 1 | 1 | 85.0 | 66.0 | 29.00000 | 155.548223 |
| 26.6 | | | | | |
| 2 | 8 | 183.0 | 64.0 | 29.15342 | 155.548223 |
| 23.3 | | | | | |

| | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627 | 50 | 1 |
| 1 | 0.351 | 31 | 0 |
| 2 | 0.672 | 32 | 1 |

#Feature scaling using MinMaxScaler

```
from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler(feature_range = (0, 1))
dataset_scaled = sc.fit_transform(data)
```

```
dataset_scaled = pd.DataFrame(dataset_scaled)
dataset_scaled
```

| | 0 | 1 | 2 | 3 | 4 | 5 |
|----------|----------|----------|----------|----------|----------|----------|
| 6 \ | | | | | | |
| 0 | 0.352941 | 0.670968 | 0.489796 | 0.304348 | 0.170130 | 0.314928 |
| 0.234415 | | | | | | |
| 1 | 0.058824 | 0.264516 | 0.428571 | 0.239130 | 0.170130 | 0.171779 |
| 0.116567 | | | | | | |
| 2 | 0.470588 | 0.896774 | 0.408163 | 0.240798 | 0.170130 | 0.104294 |
| 0.253629 | | | | | | |
| 3 | 0.058824 | 0.290323 | 0.428571 | 0.173913 | 0.096154 | 0.202454 |
| 0.038002 | | | | | | |
| 4 | 0.000000 | 0.600000 | 0.163265 | 0.304348 | 0.185096 | 0.509202 |
| 0.943638 | | | | | | |
| .. | ... | ... | ... | ... | ... | ... |
| ... | | | | | | |
| 763 | 0.588235 | 0.367742 | 0.530612 | 0.445652 | 0.199519 | 0.300613 |
| 0.039710 | | | | | | |
| 764 | 0.117647 | 0.503226 | 0.469388 | 0.217391 | 0.170130 | 0.380368 |
| 0.111870 | | | | | | |
| 765 | 0.294118 | 0.496774 | 0.489796 | 0.173913 | 0.117788 | 0.163599 |
| 0.071307 | | | | | | |
| 766 | 0.058824 | 0.529032 | 0.367347 | 0.240798 | 0.170130 | 0.243354 |
| 0.115713 | | | | | | |
| 767 | 0.058824 | 0.316129 | 0.469388 | 0.260870 | 0.170130 | 0.249489 |
| 0.101196 | | | | | | |

| | 7 | 8 |
|-----|----------|-----|
| 0 | 0.483333 | 1.0 |
| 1 | 0.166667 | 0.0 |
| 2 | 0.183333 | 1.0 |
| 3 | 0.000000 | 0.0 |
| 4 | 0.200000 | 1.0 |
| .. | ... | ... |
| 763 | 0.700000 | 0.0 |
| 764 | 0.100000 | 0.0 |
| 765 | 0.150000 | 0.0 |
| 766 | 0.433333 | 1.0 |
| 767 | 0.033333 | 0.0 |

```
[768 rows x 9 columns]
```

```
x=dataset_scaled.drop([8],axis=1)
y=dataset_scaled[[8]]
```

```
# Splitting X and Y
```

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size =
0.20, random_state = 42, stratify = data['Outcome'] )
```

```
print('The size of our train of "x" is ',x_train.shape)
print('The size of our test of "x" is ',x_test.shape)
print('The size of our train of "y" is ',y_train.shape)
print('The size of our test of "y" is ',y_test.shape)
```

```
The size of our train of "x" is (614, 8)
The size of our test of "x" is (154, 8)
The size of our train of "y" is (614, 1)
The size of our test of "y" is (154, 1)
```

```
x_train.head(3)
```

| | 0 | 1 | 2 | 3 | 4 | 5 |
|----------|----------|----------|----------|----------|----------|----------|
| 6 \ | | | | | | |
| 353 | 0.058824 | 0.296774 | 0.387755 | 0.054348 | 0.034856 | 0.184049 |
| 0.214347 | | | | | | |
| 711 | 0.294118 | 0.529032 | 0.551020 | 0.217391 | 0.009615 | 0.233129 |
| 0.154142 | | | | | | |
| 373 | 0.117647 | 0.393548 | 0.346939 | 0.358696 | 0.096154 | 0.341513 |
| 0.062767 | | | | | | |

| | 7 |
|-----|----------|
| 353 | 0.050000 |
| 711 | 0.316667 |
| 373 | 0.066667 |

```
y_train.head(3)
```

| | 8 |
|-----|-----|
| 353 | 0.0 |
| 711 | 0.0 |
| 373 | 0.0 |

```
x_test.head(3)
```

| | 0 | 1 | 2 | 3 | 4 | 5 |
|----------|----------|----------|----------|----------|----------|----------|
| 6 \ | | | | | | |
| 44 | 0.411765 | 0.741935 | 0.408163 | 0.240798 | 0.170130 | 0.188139 |
| 0.092229 | | | | | | |
| 672 | 0.588235 | 0.154839 | 0.836735 | 0.173913 | 0.042067 | 0.353783 |
| 0.088386 | | | | | | |

```
700  0.117647  0.503226  0.530612  0.217391  0.223558  0.361963
0.172929
```

```
      7
44    0.316667
672   0.433333
700   0.083333
```

```
y_test.head(3)
```

```
      8
44    0.0
672   0.0
700   0.0
```

Data Modelling

1. Logistic Regression Algorithm

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(random_state = 42)
lr.fit(x_train, y_train)
```

```
LogisticRegression(random_state=42)
```

```
print(lr.coef_)
print(lr.intercept_)
y_train_pred=pd.DataFrame(lr.predict(x_train))
y_train_pred
```

```
[[1.33192165  4.39310948  0.29586671  0.78907784  0.50191049  2.86163994
  1.15755148  0.78847976]]
[-4.90126158]
```

```
      0
0     0.0
1     0.0
2     0.0
3     0.0
4     0.0
..    ...
609   0.0
610   0.0
611   0.0
612   0.0
613   0.0
```

```
[614 rows x 1 columns]
```

```
y_test_pred=pd.DataFrame(lr.predict(x_test))
y_test_pred
```

```

      0
0      1.0
1      0.0
2      0.0
3      0.0
4      0.0
..    ...
149    0.0
150    0.0
151    0.0
152    1.0
153    0.0

```

```
[154 rows x 1 columns]
```

Accuracy score

```

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, plot_confusion_matrix
print('train accuracy:', accuracy_score(y_train_pred, y_train))
print('test accuracy:', accuracy_score(y_test_pred, y_test))

```

```
train accuracy: 0.7866449511400652
```

```
test accuracy: 0.7077922077922078
```

Confusion matrix

```

print("Train confusion_matrix \n \n ", confusion_matrix(y_train,
y_train_pred))
print("Test confusion_matrix \n \n ", confusion_matrix(y_test,
y_test_pred))

```

```
Train confusion_matrix
```

```

[[363  37]
 [ 94 120]]

```

```
Test confusion_matrix
```


```

[[83 17]
 [28 26]]

```


```
plot_confusion_matrix(lr, x_train, y_train)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x2e0e15fb520>
```



```
plot_confusion_matrix(lr,x_test,y_test)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2e0e4653fa0>
```



Classification report

```
print("Train classification_report \n \n",classification_report(y_train, y_train_pred))
```

```
print("Test classification_report \n \n",
      classification_report(y_test, y_test_pred))
```

Train classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.79 | 0.91 | 0.85 | 400 |
| 1.0 | 0.76 | 0.56 | 0.65 | 214 |
| accuracy | | | 0.79 | 614 |
| macro avg | 0.78 | 0.73 | 0.75 | 614 |
| weighted avg | 0.78 | 0.79 | 0.78 | 614 |


Test classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.75 | 0.83 | 0.79 | 100 |
| 1.0 | 0.60 | 0.48 | 0.54 | 54 |
| accuracy | | | 0.71 | 154 |
| macro avg | 0.68 | 0.66 | 0.66 | 154 |
| weighted avg | 0.70 | 0.71 | 0.70 | 154 |


ROC

Roc curve is another common tool used with binary classifiers. The dotted line represents the ROC curve of a purely random classifier; a good classifier stays as far away from that line as possible

```
from sklearn.metrics import plot_roc_curve
plot_roc_curve(lr,x_train,y_train)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_train')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



```
plot_roc_curve(lr,x_test,y_test)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



KNN

Plotting a graph for n_neighbors

```
from sklearn import metrics
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
x_axis = list(range(1, 31))
```

```
acc = pd.Series()
```

```
x = range(1,31)
```

```
for i in list(range(1, 31)):
```

```
    knn_model = KNeighborsClassifier(n_neighbors = i)
```

```
    knn_model.fit(x_train, y_train)
```

```
    prediction = knn_model.predict(x_test)
```

```
    acc = acc.append(pd.Series(metrics.accuracy_score(prediction, y_test)))
```

```
plt.plot(x_axis, acc)
```

```
plt.xticks(x)
```

```
plt.title("Finding best value for n_estimators")
```


```
plt.xlabel("n_estimators")
```

```
plt.ylabel("Accuracy")
```

```
plt.grid()
```

```
plt.show()
```

```
print('Highest value: ',acc.values.max())
```

Highest value: 0.7597402597402597

K nearest neighbors Algorithm

```
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors = 24, metric = 'minkowski', p =
2)
knn.fit(x_train, y_train)
```

```
KNeighborsClassifier(n_neighbors=24)
```

```
Knn_y_pred_train =pd.DataFrame(knn.predict(x_train))
Knn_y_pred_train
```

```

0
0    0.0
1    0.0
2    0.0
3    0.0
4    0.0
..
609  0.0
610  0.0
611  0.0
612  0.0
613  0.0
```

[614 rows x 1 columns]

```
Knn_y_pred_test =pd.DataFrame(knn.predict(x_test))
Knn_y_pred_test
```

```
      0
0    1.0
1    0.0
2    0.0
3    0.0
4    0.0
..    ..
149  0.0
150  0.0
151  0.0
152  1.0
153  0.0
```

```
[154 rows x 1 columns]
```

```
print('train accuracy:',accuracy_score(Knn_y_pred_train,y_train))
print('test accuracy:',accuracy_score(Knn_y_pred_test,y_test))
```

```
train accuracy: 0.7931596091205212
test accuracy: 0.7272727272727273
```

```
print("Train confusion_matrix \n \n
",confusion_matrix(y_train,Knn_y_pred_train ))
print("Test confusion_matrix \n \n
",confusion_matrix(y_test,Knn_y_pred_test ))
```

```
Train confusion_matrix
```


```
[[364  36]
 [ 91 123]]
```

```
Test confusion_matrix
```

```
[[87 13]
 [29 25]]
```


```
plot_confusion_matrix(knn,x_train,y_train)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x2e0dddb7b20>
```



```
plot_confusion_matrix(knn,x_test,y_test)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2e0dc8ccdf0>
```



```
print("Train classification_report \n \n",classification_report(y_train, Knn_y_pred_train))
```

Train classification_report


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.80 | 0.91 | 0.85 | 400 |
| 1.0 | 0.77 | 0.57 | 0.66 | 214 |
| accuracy | | | 0.79 | 614 |
| macro avg | 0.79 | 0.74 | 0.76 | 614 |
| weighted avg | 0.79 | 0.79 | 0.78 | 614 |

```
print("Test classification_report \n \n",classification_report(y_test, Knn_y_pred_test))
```


Test classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.75 | 0.87 | 0.81 | 100 |
| 1.0 | 0.66 | 0.46 | 0.54 | 54 |
| accuracy | | | 0.73 | 154 |
| macro avg | 0.70 | 0.67 | 0.67 | 154 |
| weighted avg | 0.72 | 0.73 | 0.71 | 154 |

```
plot_roc_curve(knn,x_train,y_train)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_train')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



```
plot_roc_curve(knn,x_test,y_test)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_train')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



Decesion tree classifier

```
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn import tree
```

```
dtc=DecisionTreeClassifier(criterion="entropy", max_depth=4, min_samples
_split= 4)
dtc.fit(x_train, y_train)
```

```
DecisionTreeClassifier(criterion='entropy', max_depth=4,
min_samples_split=4)
```

```
dtc_y_train_pred=pd.DataFrame(dtc.predict(x_train))
dtc_y_train_pred
```

```

      0
0    0.0
1    1.0
2    0.0
3    1.0
4    0.0
..    ...
609  0.0
610  0.0
611  0.0
612  1.0
613  1.0
```

```
[614 rows x 1 columns]
```

```
dtc_y_test_pred=pd.DataFrame(dtc.predict(x_test))
dtc_y_test_pred
```

```
      0
0    1.0
1    0.0
2    0.0
3    1.0
4    0.0
..    ..
149  0.0
150  0.0
151  0.0
152  1.0
153  0.0
```

```
[154 rows x 1 columns]
```

```
print('train accuracy:',accuracy_score(dtc_y_train_pred,y_train))
print('test accuracy:',accuracy_score(dtc_y_test_pred,y_test))
```

```
train accuracy: 0.7768729641693811
test accuracy: 0.7922077922077922
```

```
print("Train confusion_matrix \n \n ",confusion_matrix(y_train,
dtc_y_train_pred))
print("Test confusion_matrix \n \n ",confusion_matrix(y_test,
dtc_y_test_pred))
```

```
Train confusion_matrix
```


```
[[310  90]
 [ 47 167]]
```

```
Test confusion_matrix
```

```
[[80 20]
 [12 42]]
```


```
plot_confusion_matrix(dtc,x_train,y_train)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x2e0ddae9460>
```



```
plot_confusion_matrix(dtc,x_test,y_test)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2e0ddadff40>
```



```
print("Train classification_report \n \n",classification_report(y_train, dtc_y_train_pred))
```



```
print("Test classification_report \n \n
",classification_report(y_test, dtc_y_test_pred))
```


Train classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.87 | 0.78 | 0.82 | 400 |
| 1.0 | 0.65 | 0.78 | 0.71 | 214 |
| accuracy | | | 0.78 | 614 |
| macro avg | 0.76 | 0.78 | 0.76 | 614 |
| weighted avg | 0.79 | 0.78 | 0.78 | 614 |


Test classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.87 | 0.80 | 0.83 | 100 |
| 1.0 | 0.68 | 0.78 | 0.72 | 54 |
| accuracy | | | 0.79 | 154 |
| macro avg | 0.77 | 0.79 | 0.78 | 154 |
| weighted avg | 0.80 | 0.79 | 0.80 | 154 |

```
plot_roc_curve(dtc,x_train,y_train)
plt.plot([0, 1], [0, 1], 'g--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



```
plot_roc_curve(dtc,x_test,y_test)
plt.plot([0, 1], [0, 1], 'g--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



RandomForestClassifier

```
from sklearn.ensemble import RandomForestClassifier
RR=RandomForestClassifier(criterion='entropy',max_depth=6,n_estimators
=25)
RR.fit(x_train,y_train)
```

```
RandomForestClassifier(criterion='entropy', max_depth=6,
n_estimators=25)
```

```
RR_y_train_pred=pd.DataFrame(RR.predict(x_train))
RR_y_train_pred
```

```

      0
0    0.0
1    0.0
2    0.0
3    0.0
4    0.0
..    ..
609  0.0
610  0.0
611  0.0
612  0.0
613  0.0
```

```
[614 rows x 1 columns]
```

```
RR_y_test_pred=pd.DataFrame(RR.predict(x_test))
RR_y_test_pred
```

```
      0
0    1.0
1    0.0
2    0.0
3    0.0
4    0.0
..    ..
149  0.0
150  0.0
151  0.0
152  1.0
153  0.0
```

```
[154 rows x 1 columns]
```

```
print('train accuracy:',accuracy_score(RR_y_train_pred,y_train))
print('test accuracy:',accuracy_score(RR_y_test_pred,y_test))
```

```
train accuracy: 0.8778501628664495
test accuracy: 0.7662337662337663
```

```
print("Train confusion_matrix \n \n ",confusion_matrix(y_train,
RR_y_train_pred))
print("TEST confusion_matrix \n \n ",confusion_matrix(y_test,
RR_y_test_pred))
```

```
Train confusion_matrix
```


```
[[385  15]
 [ 60 154]]
```

```
TEST confusion_matrix
```

```
[[89 11]
 [25 29]]
```


```
plot_confusion_matrix(RR,x_train,y_train)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x2e0defe39d0>
```



```
plot_confusion_matrix(RR,x_test,y_test)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2e0def10e20>
```



```
print("Train classification_report \n \n",classification_report(y_train, RR_y_train_pred))
```

```
print("TEST classification_report \n \n
",classification_report(y_test, RR_y_test_pred))
```

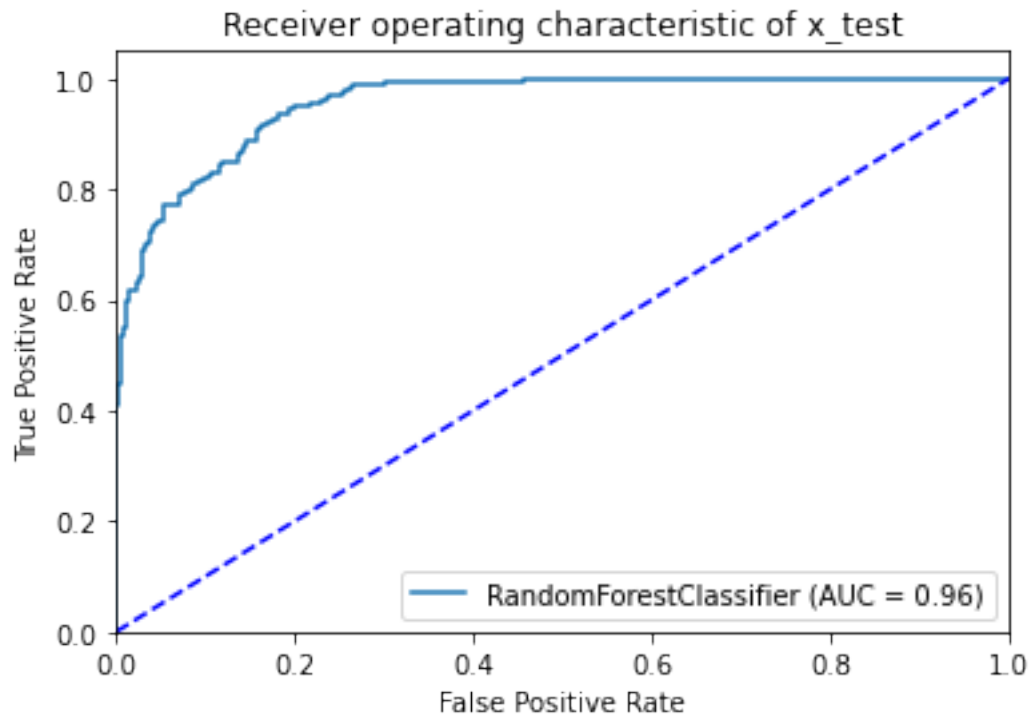
Train classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.87 | 0.96 | 0.91 | 400 |
| 1.0 | 0.91 | 0.72 | 0.80 | 214 |
| accuracy | | | 0.88 | 614 |
| macro avg | 0.89 | 0.84 | 0.86 | 614 |
| weighted avg | 0.88 | 0.88 | 0.87 | 614 |


TEST classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.78 | 0.89 | 0.83 | 100 |
| 1.0 | 0.72 | 0.54 | 0.62 | 54 |
| accuracy | | | 0.77 | 154 |
| macro avg | 0.75 | 0.71 | 0.72 | 154 |
| weighted avg | 0.76 | 0.77 | 0.76 | 154 |

```
plot_roc_curve(RR,x_train,y_train)
plt.plot([0, 1], [0, 1], 'b--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



```
plot_roc_curve(RR,x_test,y_test)
plt.plot([0, 1], [0, 1], 'b--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



AdaBoostClassifier

```
from sklearn.ensemble import AdaBoostClassifier
abc=AdaBoostClassifier(n_estimators=16)
abc.fit(x_train,y_train)
```

```
AdaBoostClassifier(n_estimators=16)
```

```
ad_pred_y_train=pd.DataFrame(abc.predict(x_train))
ad_pred_y_train
```

```
0
0    0.0
1    1.0
2    0.0
3    0.0
4    0.0
..
609  0.0
610  0.0
611  0.0
612  1.0
613  0.0
```

```
[614 rows x 1 columns]
```

```
ad_pred_y_test=pd.DataFrame(abc.predict(x_test))
ad_pred_y_test
```



```

      0
0      1.0
1      0.0
2      0.0
3      1.0
4      0.0
..    ...
149    0.0
150    0.0
151    0.0
152    1.0
153    0.0

```

```
[154 rows x 1 columns]
```

```

print('train accuracy:',accuracy_score(ad_pred_y_train,y_train))
print('test accuracy:',accuracy_score(ad_pred_y_test,y_test))

```

```

train accuracy: 0.8094462540716613
test accuracy: 0.7662337662337663

```

```

print("Train confusion_matrix \n \n ",confusion_matrix(y_train,
ad_pred_y_train))
print("Test confusion_matrix \n \n ",confusion_matrix(y_test,
ad_pred_y_test))

```

```
Train confusion_matrix
```

```

[[362  38]
 [ 79 135]]

```

```
Test confusion_matrix
```

```

[[86 14]
 [22 32]]


```

```
plot_confusion_matrix(abc,x_train,y_train)
```

```


<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x2e0e2ebe220>

```



```
plot_confusion_matrix(abc,x_test,y_test)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2e0df0a5f10>
```



```
print("Train classification_report \n \n",classification_report(y_train,ad_pred_y_train))
```

```
print("Test classification_report \n \n
",classification_report(y_test,ad_pred_y_test))
```

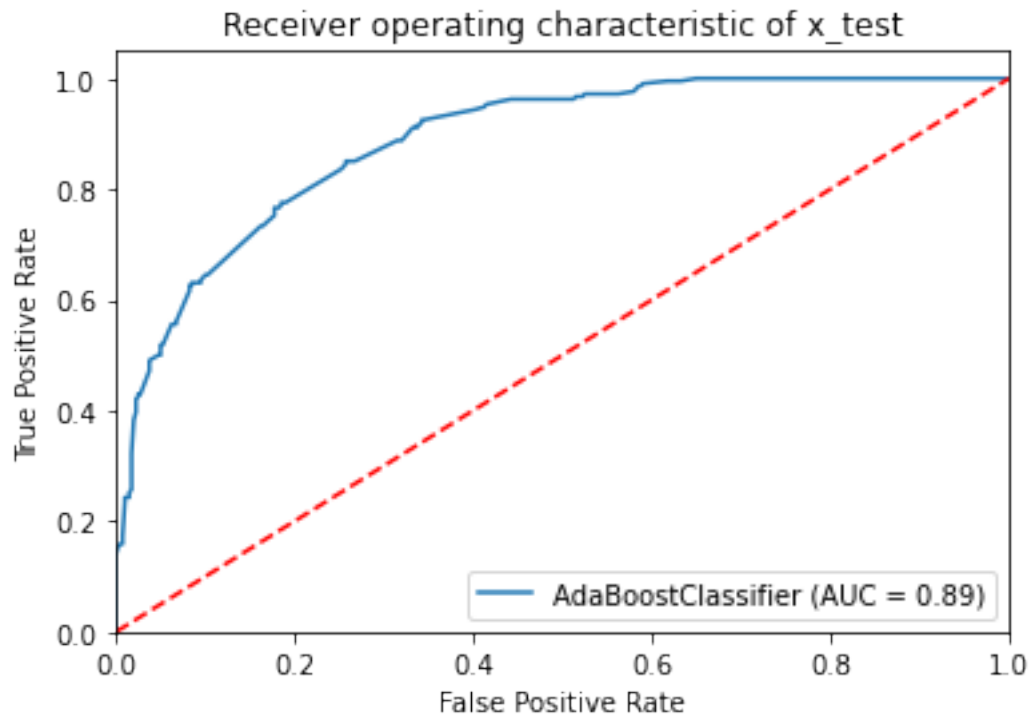
Train classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.82 | 0.91 | 0.86 | 400 |
| 1.0 | 0.78 | 0.63 | 0.70 | 214 |
| accuracy | | | 0.81 | 614 |
| macro avg | 0.80 | 0.77 | 0.78 | 614 |
| weighted avg | 0.81 | 0.81 | 0.80 | 614 |


Test classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.80 | 0.86 | 0.83 | 100 |
| 1.0 | 0.70 | 0.59 | 0.64 | 54 |
| accuracy | | | 0.77 | 154 |
| macro avg | 0.75 | 0.73 | 0.73 | 154 |
| weighted avg | 0.76 | 0.77 | 0.76 | 154 |

```
plot_roc_curve(abc,x_train,y_train)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



```
plot_roc_curve(abc,x_test,y_test)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



GradientBoostingClassifier

```
from sklearn.ensemble import GradientBoostingClassifier
gbc=GradientBoostingClassifier(max_depth=2,n_estimators= 28)
gbc.fit(x_train,y_train)
```

```
GradientBoostingClassifier(max_depth=2, n_estimators=28)
```

```
gbc_y_train_pred=pd.DataFrame(gbc.predict(x_train))
gbc_y_train_pred
```

```

      0
0    0.0
1    0.0
2    0.0
3    0.0
4    0.0
..    ..
609  0.0
610  0.0
611  0.0
612  0.0
613  0.0
```

```
[614 rows x 1 columns]
```

```
gbc_y_test_pred=pd.DataFrame(gbc.predict(x_test))
gbc_y_test_pred
```

```
      0
0      1.0
1      0.0
2      0.0
3      0.0
4      0.0
..    ...
149    0.0
150    0.0
151    0.0
152    1.0
153    0.0
```

```
[154 rows x 1 columns]
```

```
print('train accuracy:',accuracy_score(gbc_y_train_pred,y_train))
print('test accuracy:',accuracy_score(gbc_y_test_pred,y_test))
```

```
train accuracy: 0.8094462540716613
test accuracy: 0.7077922077922078
```

```
print("Train confusion_matrix \n \n ",confusion_matrix(y_train,
gbc_y_train_pred))
print("Test confusion_matrix \n \n ",confusion_matrix(y_test,
gbc_y_test_pred))
```

```
Train confusion_matrix
```


```
[[365  35]
 [ 82 132]]
```

```
Test confusion_matrix
```

```
[[83 17]
 [28 26]]
```


```
plot_confusion_matrix(gbc,x_train,y_train)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x2e0e335cb80>
```



```
plot_confusion_matrix(gbc,x_test,y_test)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at  
0x2e0ddb85e50>
```



```
print("Train classification_report \n \n  
",classification_report(y_train,gbc_y_train_pred))
```

```
print("Test classification_report \n \n",
      classification_report(y_test,gbc_y_test_pred))
```


Train classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.82 | 0.91 | 0.86 | 400 |
| 1.0 | 0.79 | 0.62 | 0.69 | 214 |
| accuracy | | | 0.81 | 614 |
| macro avg | 0.80 | 0.76 | 0.78 | 614 |
| weighted avg | 0.81 | 0.81 | 0.80 | 614 |


Test classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.75 | 0.83 | 0.79 | 100 |
| 1.0 | 0.60 | 0.48 | 0.54 | 54 |
| accuracy | | | 0.71 | 154 |
| macro avg | 0.68 | 0.66 | 0.66 | 154 |
| weighted avg | 0.70 | 0.71 | 0.70 | 154 |

```
plot_roc_curve(gbc,x_train,y_train)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```

```
plot_roc_curve(gbc,x_test,y_test)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



SVC

Support Vector Classifier Algorithm

```
from sklearn.svm import SVC
svc = SVC(kernel = 'linear', random_state = 42)
svc.fit(x_train, y_train)
```

```
SVC(kernel='linear', random_state=42)
```

```
svc_y_train_pred=pd.DataFrame(svc.predict(x_train))
svc_y_train_pred
```

```
0
0  0.0
1  0.0
2  0.0
3  0.0
4  0.0
..
609 0.0
610 0.0
611 0.0
612 0.0
613 0.0
```

```
[614 rows x 1 columns]
```

```
svc_y_test_pred=pd.DataFrame(svc.predict(x_test))
svc_y_test_pred
```

```
      0
0    1.0
1    0.0
2    0.0
3    0.0
4    0.0
..    ..
149  0.0
150  0.0
151  0.0
152  1.0
153  0.0
```

```
[154 rows x 1 columns]
```

```
print('train accuracy:',accuracy_score(svc_y_train_pred,y_train))
print('test accuracy:',accuracy_score(svc_y_test_pred,y_test))
```

```
train accuracy: 0.7817589576547231
test accuracy: 0.7012987012987013
```

```
print("Train confusion_matrix \n \n ",confusion_matrix(y_train,
svc_y_train_pred))
print("Test confusion_matrix \n \n ",confusion_matrix(y_test,
svc_y_test_pred))
```

```
Train confusion_matrix
```


```
[[363  37]
 [ 97 117]]
```

```
Test confusion_matrix
```

```
[[83 17]
 [29 25]]
```


```
plot_confusion_matrix(svc,x_train,y_train)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x2e0ddd6bbb0>
```



```
plot_confusion_matrix(svc,x_test,y_test)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2e0e335ce80>
```



```
print("Train classification_report \n \n",classification_report(y_train,svc_y_train_pred))
```

```
print("Test classification_report \n \n
",classification_report(y_test,svc_y_test_pred))
```


Train classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.79 | 0.91 | 0.84 | 400 |
| 1.0 | 0.76 | 0.55 | 0.64 | 214 |
| accuracy | | | 0.78 | 614 |
| macro avg | 0.77 | 0.73 | 0.74 | 614 |
| weighted avg | 0.78 | 0.78 | 0.77 | 614 |


Test classification_report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.74 | 0.83 | 0.78 | 100 |
| 1.0 | 0.60 | 0.46 | 0.52 | 54 |
| accuracy | | | 0.70 | 154 |
| macro avg | 0.67 | 0.65 | 0.65 | 154 |
| weighted avg | 0.69 | 0.70 | 0.69 | 154 |

```
plot_roc_curve(svc,x_train,y_train)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



```
plot_roc_curve(svc,x_test,y_test)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic of x_test')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```



Model Evaluation

data.columns

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
       'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

```
models = []
models.append(('Logistic Regression', lr))
models.append(('KNN', knn))
models.append(('SVC', svc))
models.append(('Decision tree', dtc))
models.append(('Random Forest', RF))
models.append(('Adboost', abc))
models.append(('Gboost', gbc))
```

```
from sklearn import model_selection
model_names=['Logistic Regression','KNN','SVC','Decision tree','Random
Forest','Adboost','Gboost']
results = []
reports = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']
scoring = 'accuracy'
for report, model in models:
    kfold = model_selection.KFold(n_splits=10, random_state=7)
    cv_results = model_selection.cross_val_score(model, x_train,
```

```

y_train , cv=kfold, scoring=scoring)
    results.append(cv_results)
    reports.append(report)
    msg = "%s: %f (%f)" % (report,cv_results.mean(), cv_results.std())
    print(msg)
# boxplot algorithm comparison
plt.suptitle('Algorithm Comparison')
plt.boxplot(results)
plt.xlabel('models')
plt.show()

Logistic Regression: 0.781756 (0.038181)
KNN: 0.770307 (0.044869)
SVC: 0.776811 (0.043028)
Decision tree: 0.744209 (0.060441)
Random Forest: 0.752353 (0.034431)
Adboost: 0.745796 (0.054192)
Gboost: 0.773533 (0.048587)

```

